

CS1010 Laboratory 09

Backtracking

Zhang Puyu

Group BD04

November 7, 2024

Plan of the Day

CS1010

Laboratory 09

Zhang Puyu

Exercise 6
Review

Recursion &
Backtracking

Graph Traversal

Exercise 7
Review

1 Exercise 6 Review

2 Recursion & Backtracking
■ Graph Traversal

3 Exercise 7 Review

sort.c: Two Common Approaches

CS1010

Laboratory 09

Zhang Puyu

Exercise 6
Review

Recursion &
Backtracking
Graph Traversal

Exercise 7
Review

Approach 1: Small to Big.

- 1 Locate the index of the minimum element using either linear or binary search.
- 2 Put this minimum as the first element of output array.
- 3 Consider the left and right neighbours of the minimum:
 - If `left > right`, we put `right` into the next index of output array, and move one index to the right.
 - Otherwise, put `left` into the next index of output array, and move one index to the left.
- 4 When one side reaches the end, put whatever is left in the other side into the output array in ascending order.

sort.c: Two Common Approaches

CS1010

Laboratory 09

Zhang Puyu

Exercise 6
Review

Recursion &
Backtracking
Graph Traversal

Exercise 7
Review

Approach 2: Big to Small.

- 1 Initialise `left` at index 0 and `right` at index `len - 1`.
 - If `left > right`, we put `left` into the next index of output array, and move one index to the right.
 - Otherwise, put `right` into the next index of output array, and move one index to the left.
- 2 We stop when `left > right`. The output list now is reversely sorted.
- 3 Since we know the size of the output array, we can easily print it in reverse order.

- 1 We will look for the minimum between `start` and `end`.
- 2 Compute the midpoint index `mid`.
 - If `mid + 1` is a bigger element, then we continue searching between `start` and `mid`.
 - Otherwise, we will search in between `mid + 1` and `end`.
- 3 The above terminates when the search space only contains a singleton.

inversion.c: How to Count Cumulatively

CS1010

Laboratory 09

Zhang Puyu

Exercise 6
Review

Recursion &
Backtracking
Graph Traversal

Exercise 7
Review

- 1 Let $i = 0$ and $j = 0$. Assume the array is of size n .
- 2 For the i -th element, do:
 - If $\text{arr}[n - 1 - j] < \text{arr}[i]$, increment j .
 - Repeat until $\text{arr}[n - 1 - j] \geq \text{arr}[i]$.
 - j is the number of elements which are inversions with respect to $\text{arr}[i]$.
 - Increment total count by j and move to the next element.
- 3 When $i = n - 1 - j$, we claim that $\text{arr}[i]$ is the peak.
- 4 From $\text{arr}[i]$ to $\text{arr}[n - 1]$ there are $\frac{(n-i)(n-i-1)}{2}$ inversions.

Question: Why do we not reset j ?

mark.c: Stable Sort

CS1010

Laboratory 09

Zhang Puyu

Exercise 6
Review

Recursion &
Backtracking
Graph Traversal

Exercise 7
Review

- Note that equal marks are already sorted based on their names in the input.
- So we only need to run counting sort on the marks, with the extra constraint that:
- The names appear in the exact same order as in the input.
- This is known as a **stable** sorting algorithm.

Definition (Stable Sort)

A sorting algorithm is **stable** if for any $a = b$ in the input such that a comes before b , in the sorted output a still comes before b .

mark.c: Stable Sort

CS1010

Laboratory 09

Zhang Puyu

Exercise 6
Review

Recursion &
Backtracking
Graph Traversal

Exercise 7
Review

- We can compute the frequency table of all marks.

...	m_i	m_{i+1}	m_{i+2}	m_{i+3}	m_{i+4}	...
...	f_i	f_{i+1}	f_{i+2}	f_{i+3}	f_{i+4}	...

- Simpler problem: how to transform the frequency table into a cumulative frequency table?

$$F_i = \sum_{j=0}^i f_j.$$

- The **cumulative frequency** table:

...	m_i	m_{i+1}	...
...	$F_i = F_{i-1} + f_i$	$F_{i+1} = F_i + f_{i+1}$...

- Claim: If the cumulative frequency of m_k is F_k , then in the sorted array, the n -th occurrence of m_k is at index $F_k - f_k + n - 1$.

Recursion & Backtracking

CS1010

Laboratory 09

Zhang Puyu

Exercise 6
Review

Recursion &
Backtracking

Graph Traversal

Exercise 7
Review

- Previously, we have seen how to use recursion to arrive at $P(k)$ given $P(1)$ and $P(n) \mapsto P(n+1)$.
- It seems (erroneously) that recursion is always about “going backwards”, but it is not always the case.
- A problem-solving paradigm:
 - 1 Suppose we have an **initial state** S .
 - 2 Our **goal** is to find all states where a **terminating condition** p holds.
 - 3 Upon entering each **current state** X , we first check if p is true. If it is, we **collect** X as a valid result.
 - 4 Else, we explore the **actions** that can be performed at state X to **transition** to a different state Y .
 - 5 For each new state Y generated this way, we enter it and repeat the process.
 - 6 After **exhausting** all possible actions, we **return** to the previous state.
- **Backtracking** is searching by brute force.

An Example: `stone.c` from past year PE

CS1010

Laboratory 09

Zhang Puyu

Exercise 6
Review

Recursion &
Backtracking

Graph Traversal

Exercise 7
Review

Basically, this asks you to **print out all binary sequences of size n in ascending order**.

- 1 We start with our sequence $s = \emptyset$ and number of digits $k = 0$.
- 2 At every stage, check: is $k = n$?
 - If yes, then this is a valid sequence, we will print it and return. (**Collect** a valid result and **backtrack** to the previous state.)
 - If no, then we have 2 choices to proceed to the next stage:
 - 1 Append 0 to the end of s ; or
 - 2 Append 1 to the end of s .

An Example: stone.c from past year PE

CS1010

Laboratory 09

Zhang Puyu

Exercise 6
Review

Recursion &
Backtracking

Graph Traversal

Exercise 7
Review

```
void print_binary(long k, long n) {
    // Is [s, k] a goal state?
    if (k == n) {
        // Yes, collect the result.
        cs1010_println_string("");
        // Backtrack to the previous state.
        return;
    }
    // No, explore the possible actions.
    // Action 1: append 0.
    // Transition 1: [s, k] -> [s0, k + 1].
    putchar('0');
    print_binary(k + 1, n);
    // Action 2: append 1.
    // Transition 2: [s, k] -> [s1, k + 1].
    putchar('1');
    print_binary(k + 1, n);
}

int main() {
    long n = cs1010_read_long();
    print_binary(0, n); // Initial state: ["", 0].
    return 0;
}
```

Graph Traversal: Depth-First Search

CS1010

Laboratory 09

Zhang Puyu

Exercise 6
Review

Recursion &
Backtracking
Graph Traversal

Exercise 7
Review

A graph consists of vertices and edges. An edge $x \rightarrow y$ allows you to move from vertex x to vertex y .

- Given two vertices v and u , we wish to determine whether they are connected by a path.
- Identify the following:
 - 1 **Current Stage:** The current vertex x (initially v).
 - 2 **Terminating Condition:** $x = u$.
 - 3 **State Transitions:** Move to any unvisited neighbour of x .
- DFS at VisuAlgo.

Other Interesting Applications of Backtracking

CS1010

Laboratory 09

Zhang Puyu

Exercise 6
Review

Recursion &
Backtracking
Graph Traversal

Exercise 7
Review

- Writing a parser and compiler.
- Goal-Oriented Action Planning in game AI.
- Generating permutations.
- Building a dependency tree.

- This is a graph traversal problem in a grid. So each vertex can be uniquely represented by lattice point coordinates.
- Identify the following:
 - 1 Current Stage:** The current coordinates (a, b) (initially $(0, 0)$).
 - 2 Terminating Condition:** $(a, b) = (x, y)$.
 - 3 State Transitions:**
 - 1** $(a, b) \mapsto (a + 1, b)$
 - 2** $(a, b) \mapsto (a, b + 1)$

In fact this is a **troll question...** The phrasing seems that $\mathcal{O}(xy)$ is optimal, but someone trained in combinatorics quickly realises that we can do this in $\mathcal{O}(\min\{x, y\})$ (and using a single loop only).

walk.c: The Uninteresting Solution

CS1010

Laboratory 09


Zhang Puyu

Exercise 6
Review

Recursion &
Backtracking
Graph Traversal

Exercise 7
Review

How to go from A to B (x steps to the right and y steps to the top)?



Let $P(x, y)$ be the number of paths from A to (x, y) . Observe that

$$P(x, y) = \begin{cases} 1 & \text{if } x = 0 \text{ or } y = 0 \\ P(x-1, y) + P(x, y-1) & \text{otherwise} \end{cases}.$$

walk.c: The Uninteresting Solution

CS1010

Laboratory 09

Zhang Puyu

Exercise 6
Review

Recursion &
Backtracking
Graph Traversal

Exercise 7
Review

```
long move_to(long m, long n) {  
    if (m == 0 || n == 0) {  
        return 1;  
    }  
    return move_to(m - 1, n) +  
           move_to(m, n - 1);  
}
```

Have you spotted the duplicated computation here?

walk.c: The Uninteresting Solution

CS1010


Laboratory 09

Zhang Puyu

Exercise 6
Review

Recursion &
Backtracking
Graph Traversal

Exercise 7
Review



Note that the above grid can be seen as a **2D array**! Each vertex is an entry, in which we will store the number of ways to get to that vertex from A.

Now we can use this 2D array as a look-up table to build a bottom-up solution. (So this is what "think recursively but solve iteratively" supposed to mean...)

We will need to fill up the matrix completely before we know the value of $\text{mat}[m][n]$, so it is $\mathcal{O}(mn)$.

walk.c: The Legit Solution

CS1010

Laboratory 09

Zhang Puyu

Exercise 6
Review

Recursion &
Backtracking
Graph Traversal

Exercise 7
Review

Suppose that 0 represents a step rightwards and 1 represents a step upwards, then every path can be uniquely mapped to a binary sequence with m 0's and n 1's. (Bijection)

Now we are dealing with a simpler question: **how many distinct binary sequences are there with m 0's and n 1's?**

The way to create such a sequence: **list down $(m + n)$ 0's and toggle n of them to be 1.**

Using some maths, we know that the answer is

$$\binom{m+n}{n} = \frac{(m+n)(m+n-1)(m+n-2) \cdots (m+1)}{1 \cdot 2 \cdot 3 \cdots n}.$$

walk.c: The Legit Solution

CS1010

Laboratory 09

Zhang Puyu

Exercise 6
Review

Recursion &
Backtracking
Graph Traversal

Exercise 7
Review

```
long x_choose_y(long x, long y) {
    long top = 1;
    long bottom = 1;
    for (long i = 0; i < y; i += 1) {
        top *= (x - i);
        bottom *= (i + 1);
    }
    return top / bottom; // Guaranteed to be an integer
}

int main() {
    long m = cs1010_read_long();
    long n = cs1010_read_long();
    cs1010_println_long(x_choose_y(m + n, m > n ? n : m));
}
```

You are wrong if you think this is correct. (Why?)

Both `top` and `bottom` are factorials, which easily gets too big for `long` or even `long long`.

walk.c: The Legit Solution

CS1010

Laboratory 09

Zhang Puyu

Exercise 6
Review

Recursion &
Backtracking
Graph Traversal

Exercise 7
Review

Observe that

$$\binom{x}{y} = \frac{x - y + 1}{y} \binom{x}{y - 1}$$

This implies that we can start from $c = \binom{x}{1} = x$, at each iteration, we update by $c = c * (x - i) / (i + 1)$. This number will be small at each iteration, so we can use `long` again!

walk.c: The Legit Solution

CS1010

Laboratory 09

Zhang Puyu

Exercise 6

Review

Recursion &
Backtracking

Graph Traversal

Exercise 7

Review

```
long x_choose_y(long x, long y) {
    long c = x; // x choose 1
    for (long i = 1; i < y; i += 1) {
        // Guaranteed to be an integer
        c = c * (x - i) / (i + 1);
    }
    return c;
}

int main() {
    long m = cs1010_read_long();
    long n = cs1010_read_long();
    cs1010_println_long(x_choose_y(m + n, m > n ? n : m));
}
```

This will be $\mathcal{O}(\min\{m, n\})$ (WAY faster than proposed “optimal” solution).

maze.c: A First Encounter with DFS

CS1010

Laboratory 09

Zhang Puyu

Exercise 6
Review

Recursion &
Backtracking
Graph Traversal

Exercise 7
Review

- This is a real graph traversal problem which does require recursion and DFS.
- Identify the following:
 - 1 **Current Stage:** The current coordinates (a, b) (initially at player spawn point).
 - 2 **Terminating Condition:** (a, b) is at the boundary.
 - 3 **State Transitions:** Go to any **unvisited** neighbouring point which is not a wall.

maze.c: A First Encounter with DFS

CS1010

Laboratory 09

Zhang Puyu

Exercise 6
Review

Recursion &
Backtracking
Graph Traversal

Exercise 7
Review

We are given a maze (a 2D array) and a birth point at (x, y) .

- We wish to know if we can reach the exit from (x, y) .
- Step 1: Check **whether (x, y) is already an exit.**
- Step 2: If we are not at the exit yet, we iterate the neighbouring cells. For each neighbour (x', y') :
 - If (x', y') is visited, skip.
 - We can use another 2D array (the “map”) to store the status of the cells!
 - Note that we actually waste a lot of memory space here. (A more space-efficient solution requires a hash table.)
 - Else, set starting point to (x', y') and run the same processes.
- Step 3: If we cannot get to the exit from any neighbouring cell, it means the maze is unsolvable.
- How do we keep track of step count? **Use a pointer** so that we can update its value in any function call!

maze.c: A First Encounter with DFS

CS1010

Laboratory 09

Zhang Puyu

Exercise 6

Review

Recursion &
Backtracking

Graph Traversal

Exercise 7

Review

```
bool can_escape(char **maze, bool **visited, long x, long y,
                long nrows, long ncols, long *count) {
    visited[x][y] = true; // Current cell becomes visited
    // TODO: Implement is_exit
    if (is_exit(x, y, nrows, ncols)) {
        return true; // Escaped!
    }
    // Try going up, TODO: Implement can_go_there and move_to
    if (can_go_there(x - 1, y, maze, visited)) {
        // Go there
        move_to(maze, x - 1, y, x, y, count, nrows);
        if (can_escape(maze, visited, x - 1, y,
                      nrows, ncols, count)) {
            return true; // That cell is connected to an exit!
        }
        // Think: What does it mean if we ever reach this line?
        move_to(maze, x, y, x - 1, y, count, nrows); // Step back
    }
    // TODO: go right, go down, and go left
    return false; // No exit is reachable
}
```


fill.c: A Maze with No Walls

CS1010

Laboratory 09

Zhang Puyu

Exercise 6
Review

Recursion &
Backtracking
Graph Traversal

Exercise 7
Review

The key to solving this question is to be able to identify it as an **implicit graph traversal question!**

- The coordinates at which we drop the colour bucket is essentially our player spawn point.
- Instead of finding an exit, this time our task is to **visit every reachable cell with the same colour as the spawn point.**
- Note that as compared the `maze.c`, now each cell contains additional information: a 3-tuple (r, g, b) indicating its colour!
 - This is a 3D array! How do we process it?
 - Approach 1: Fake 3D array by `cols = 3 * n`. To visit the cell at row x and column y , use `canva[x][y * 3]`.
 - Approach 2: Use a real 3D array `long (**canva)[3]`.
 - Approach 3: Use a real 3D array `long **canva[3]`.

fill.c: A Maze with No Walls

CS1010

Laboratory 09

Zhang Puyu

Exercise 6
Review

Recursion &
Backtracking
Graph Traversal

Exercise 7
Review

Pseudo-code:

- We wish to pour the colour bucket containing (R', G', B') at (x, y) which has colour (R, G, B) .
- If $(R, G, B) = (R', G', B')$, there's nothing to colourise.
- Else, we will change the colour of (x, y)
- For the four neighbours:
 - If the neighbour is out of boundary of the canva, skip.
 - If the neighbour has a different colour from (R, G, B) , skip.
 - Else, spread the colour (R', G', B') to the neighbour.

fill.c: A Maze with No Walls

CS1010

Laboratory 09

Zhang Puyu

Exercise 6
Review

Recursion &
Backtracking
Graph Traversal

Exercise 7
Review

```
void colour(
    long (**canva)[3], long x, long y,
    long old_colour[3], long new_colour[3],
    long nrows, long ncols
) {
    if (!is_within_boundary(x, y, nrows, ncols) ||
        are_the_same_colour(old_colour, new_colour)) {
        return;
    } // TODO: Implement these two boolean functions

    // TODO: Change the colour of current pixel
    change_colour(canva, x, y, old_colour, new_colour);

    // Now go up
    colour(canva, x - 1, y, old_colour, new_colour,
           nrows, ncols);

    // TODO: Implement go-right, go-down and go-left cases
}
```

sudoku.c: Intimidating, but the Idea Is Easy

CS1010

Laboratory 09

Zhang Puyu

Exercise 6
Review

Recursion &
Backtracking
Graph Traversal

Exercise 7
Review

As human beings, how do we solve a Sudoku puzzle?

By **trial and error!** We randomly put a number into the blank and proceed, if it works out we have a solution, else we erase the number and try a different one.

We only need to mimic this behaviour with code!

(By the way, the wording of the question statement was very ambiguous and I am personally not a fan of this problem.)

sudoku.c: Intimidating, but the Idea Is Easy

CS1010

Laboratory 09

Zhang Puyu

Exercise 6
Review

Recursion &
Backtracking
Graph Traversal

Exercise 7
Review

- We will run through the puzzle from top to bottom and left to right.
- At the x -th row and y -th column, print the current state of the puzzle and check:
 - If $x > \text{NROWS}$, we have successfully run through the puzzle and nothing wrong happens, so the puzzle is solvable.
 - Else if $y > \text{NCOLS}$, we should move to the next row.
 - Else if (x, y) is not empty, we should move to the next column.
 - Else if (x, y) is empty, for $i = 1, 2, \dots, 9$:
 - If placing i here is safe, set $(x, y) = i$ and move on.
 - If a solution can be found from here, the puzzle is solvable.
 - Else, reset (x, y) and try the next valid i .
 - If none of $1, 2, \dots, 9$ leads to a solution, the puzzle is not solvable with the current state.

sudoku.c: Anything Wrong Here?

CS1010
Laboratory 09

Zhang Puyu

Exercise 6
Review

Recursion &
Backtracking
Graph Traversal

Exercise 7
Review

```
bool can_solve(char puzzle[9][9], long x, long y) {
    print_sudoku(puzzle); // Print current state of the puzzle
    if (x > 8) {
        // Have checked every cell and nothing went wrong
        return true;
    }
    if (y > 8) {
        return can_solve(puzzle, x + 1, 0); // To next row
    }
    if (puzzle[x][y] != EMPTY) {
        return can_solve(puzzle, x, y + 1); // To next column
    }
    for (char guess = '1'; guess <= '9'; guess += 1) {
        // TODO: Implement is_valid
        if (is_valid(guess, x, y, puzzle)) {
            puzzle[x][y] = guess;
            if (can_solve(puzzle, x, y + 1)) { return true; }
            puzzle[x][y] = EMPTY; // Didn't work, guess again
        }
    }
    return false; // Tried everything, didn't find a solution
}
```